

Clustering

Tony Ni, Antonella Basso, Jose Lopez

6/21/2020

Libraries

Reading in Data

```
setwd("~/harvard-summer-biostats")
df <- read_csv("data/wide_illinois.csv") #read in data
```

```
## Parsed with column specification:
## cols(
##   .default = col_double(),
##   well_id = col_character(),
##   site = col_character(),
##   disposal_area = col_character(),
##   type = col_character(),
##   gradient = col_character()
## )

## See spec(...) for full column specifications.
```

```
df1 <- na.omit(df) #get rid of na's
```

Hierarchical Methods

If we want to look for cereal groups via hierarchical clustering, we need to construct a distance matrix. Distances are constructed with the *dist* function, and we need to choose whether we compute them on scaled or unscaled variables (standardize or not).

```
df.dist <- dist(df1[, -c(1:5)])
```

Now we look at how hierarchical clustering is applied. The relevant function is *hclust*.


```
hcsingle <- hclust(df.dist, method = "single")
list(hcsingle) # reminds you of properties of the solution, if desired
```

```
## [[1]]
##
## Call:
```

```
## hclust(d = df.dist, method = "single")
##
## Cluster method   : single
## Distance         : euclidean
## Number of objects: 191
```

This creates the solution, and we can look at the dendrogram as:

```
plot(hcsingle)
```



The options for hclust in terms of linkages are provided in the help under options for method. The following options are listed: “ward.D”, “ward.D2”, “single”, “complete”, “average”, “mcquitty”, “median” or “centroid”.

In order to obtain cluster labels, we need to *cut* our dendrograms.


```
singleSol <- (cutree(hcsingle, k = 3)) #cluster labels are numeric, k= # clusters
summary(as.factor(singleSol)) #as factor to get table
```

```
##   1   2   3
## 172 18   1
```

To learn more details about the clusters we found:

```
hclward <- hclust(df.dist, method = "ward.D")
plot(hclward, labels = df1$gradient, cex = 0.7) #cex adjusts size of label
```

Cluster Dendrogram




```
wardSol <- (cutree(hclward, k = 2)) #cluster labels are numeric, k= # clusters
summary(as.factor(wardSol)) #as factor to get table
```

```
##    1    2
## 172   19
```

```
favstats(Antimony ~ wardSol, data = df1) #can choose any variable
```

```
##   wardSol      min    Q1 median    Q3      max      mean      sd    n
## 1       1 0.0007625 0.001  0.003 0.003 0.1137778 0.002903086 0.00858458 172
## 2       2 0.5005000 1.000  1.000 1.000 1.1000000 0.926447368 0.19071691  19
##   missing
## 1         0
## 2         0
```

```
bwplot(Antimony ~ as.factor(wardSol), data = df1)
```



Our cluster sizes are extremely uneven... Our first cluster has 172 wells and the second has 19.


```
favstats(#chemical here ~ wardSol, data = df1)
bwplot(#chemical here ~ as.factor(wardSol), data = df1)
```

We can view the solution in the PC space (say 2-D) to see how well-separated the clusters are in that space. Because we used an unstandardized distance, we will run the PCA on the covariance matrix.

```
dfPCA <- princomp(df1[, -c(1:5)], cor = FALSE)

plot(dfPCA$scores[, 1:2], type = "n", xlab = "PC1", ylab = "PC2", main = "Ward's cluster solution") #bl
text(dfPCA$scores[, 1:2], labels = wardSol, cex = 0.6) #add the text
```

Ward's cluster solution



We may want to go into investigation to see what sort of traits/attributes are shared by the wells in each cluster and seeing if we can find meaning in them.

K-means Methods

For k-means, we don't need to compute the distance matrix ourselves. We feed the function the data set to operate on:

```
Ksol1 <- kmeans(scale(df1[, -c(1:5)]), centers = 2) #centers is the # of clusters desired
list(Ksol1) #so you can see what it gives you
```

```
## [[1]]
## K-means clustering with 2 clusters of sizes 16, 175
##
## Cluster means:
##      Antimony      Arsenic      Barium  Beryllium      Boron      Cadmium      Calcium
## 1  3.2161104  2.4960121  2.6978935  3.1156481  3.1096817  3.1138065  2.8270214
## 2 -0.2940444 -0.2282068 -0.2466646 -0.2848593 -0.2843138 -0.2846909 -0.2584705
##      Chloride      Chromium      Cobalt      Fluoride      Lead      Lithium      Mercury
## 1 -0.39474887  1.4858331  1.680383 -0.13626181  1.4040106  2.817204  3.2148635
## 2  0.03609133 -0.1358476 -0.153635  0.01245822 -0.1283667 -0.257573 -0.2939304
##      Molybdenum      pH Radium 226+228      Selenium      Sulfate      Thallium
## 1  2.5288750  0.58124207  0.33270030  3.1317800 -0.53202903  3.202237
## 2 -0.2312114 -0.05314213 -0.03041831 -0.2863342  0.04864265 -0.292776
##      Total Dissolved Solids
```

```
## 1          -0.46565729
## 2          0.04257438
##
## Clustering vector:
## [1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
## [38] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 2 1 2 2 1 1 1 2 2 2 2
## [75] 2 2 1 2 2 2 2 2 1 1 2 1 1 2 2 2 2 2 2 2 1 2 2 2 2 2 2 2 2 2 2 2
## [112] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
## [149] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1
## [186] 2 1 2 2 2 2
##
## Within cluster sum of squares by cluster:
## [1] 792.9099 1210.3005
## (between_SS / total_SS = 49.8 %)
##
## Available components:
##
## [1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
## [6] "betweenss"    "size"         "iter"         "ifault"
```

The list option provides us with lots of information. We can pull out the cluster means as:

```
Ksol1$centers
```


```
##      Antimony      Arsenic      Barium  Beryllium      Boron      Cadmium      Calcium
## 1  3.2161104  2.4960121  2.6978935  3.1156481  3.1096817  3.1138065  2.8270214
## 2 -0.2940444 -0.2282068 -0.2466646 -0.2848593 -0.2843138 -0.2846909 -0.2584705
##      Chloride      Chromium      Cobalt      Fluoride      Lead      Lithium      Mercury
## 1 -0.39474887  1.4858331  1.680383 -0.13626181  1.4040106  2.817204  3.2148635
## 2  0.03609133 -0.1358476 -0.153635  0.01245822 -0.1283667 -0.257573 -0.2939304
##      Molybdenum      pH Radium 226+228      Selenium      Sulfate      Thallium
## 1  2.5288750  0.58124207  0.33270030  3.1317800 -0.53202903  3.202237
## 2 -0.2312114 -0.05314213 -0.03041831 -0.2863342  0.04864265 -0.292776
##      Total Dissolved Solids
## 1          -0.46565729
## 2          0.04257438
```

In order to determine if we have chosen a “good” value of the number of clusters, we can look at the within cluster sum of squares for this solution and a few other options for k, the number of clusters. This runs the solution from 1 to 10 clusters and pulls the within group sum of squares from each.

```
n <- nrow(df1) #number of observations

wss <- rep(0, 10) #creates 10 copies of 0 to create an empty vector
for(i in 1:10){
  wss[i] <- sum(kmeans(scale(df1[, -c(1:5)]), centers = i)$withinss)
}

plot(1:10, wss, type = "b", xlab = "Number of groups", ylab = "Within groups sum of squares")
```



We look for elbows in the plot - here there are elbows at 2 and 6 (ish?), maybe these values will be good to use?

With two clusters, we should see if there is any relationship with sites...

```
tally(Ksol1$cluster ~ type, data = df1, format = "count")
```

```
##           type
## Ksol1$cluster  L  M  SI
##           1  16  0   0
##           2  32  5 138
```

There does seem to be something of interest here... The first cluster has a varied mix between all types with the majority in SI, while wells in cluster 2 seem to only consist of L.

We can compare clustering solutions with similar tables. How do the K-means and Ward's solutions overlap?

```
tally(Ksol1$cluster ~ wardSol, data = df1, format = "count")
```

```
##           wardSol
## Ksol1$cluster  1  2
##           1   0 16
##           2 172  3
```

They seem to match up fairly well!

Alternative Approaches


“If you haven’t already, you can try some alternative filtering approaches: there might be a lot of contaminants that don’t vary that much among the wells, so one thing you could do is find the contaminants with the highest between-well variability. For each chemical you can compute the variance across wells within a site, and take only the some number of contaminants (10? 20? Or 30?) with the highest variance, and then use those chemicals to do the PCA/k-means/hierarchical clustering.”

```
df2 <- df1 %>%
  gather(key, value, c(6:26)) %>%
  group_by(site, key) %>%
  mutate(m = mean(value), sd = sd(value)) %>%
  arrange(desc(sd)) %>%
  print
```

```
## # A tibble: 4,011 x 9
## # Groups:   site, key [357]
##   well_id site      disposal_area type gradient key value m sd
##   <chr> <chr>      <chr>      <chr> <chr> <chr> <dbl> <dbl> <dbl>
## 1 G03D Prairie St~ Near Field Lan~ L Upgradi~ Calc~ 5.85e4 92542. 52347.
## 2 G04D Prairie St~ Near Field Lan~ L Downgra~ Calc~ 1.03e5 92542. 52347.
## 3 G05D Prairie St~ Near Field Lan~ L Downgra~ Calc~ 7.53e4 92542. 52347.
## 4 G07D Prairie St~ Near Field Lan~ L Downgra~ Calc~ 9.43e4 92542. 52347.
## 5 G08D Prairie St~ Near Field Lan~ L Downgra~ Calc~ 7.10e4 92542. 52347.
## 6 G09D Prairie St~ Near Field Lan~ L Downgra~ Calc~ 2.77e5 92542. 52347.
## 7 G10D Prairie St~ Near Field Lan~ L Downgra~ Calc~ 1.03e5 92542. 52347.
## 8 G13D Prairie St~ Near Field Lan~ L Downgra~ Calc~ 9.51e4 92542. 52347.
## 9 G15D Prairie St~ Near Field Lan~ L Downgra~ Calc~ 1.16e5 92542. 52347.
## 10 G17D Prairie St~ Near Field Lan~ L Downgra~ Calc~ 8.94e4 92542. 52347.
## # ... with 4,001 more rows
```

```
df2 %>%
  ggplot(.) +
  theme_bw() +
  aes(x = site, y = m, ymin = m - sd, ymax = m + sd, color = key) +
  geom_point(position = position_dodge(width = 0.5)) +
  geom_errorbar(position = position_dodge(width = 0.5)) +
  ylab("Concentration") +
  ggtitle("Variability between wells among sites by contaminant") +
  coord_flip() +
  theme(legend.position="bottom")
```

```
## Warning: Removed 21 rows containing missing values (geom_errorbar).
```

Calcium dominates the entire graph... Remove it...

```
df2 %>%
  filter(key != "Calcium") %>%
  ggplot(.) +
  theme_bw() +
```

```
aes(x = site, y = m, ymin = m - sd, ymax = m + sd, color = key) +  
geom_point(position = position_dodge(width = 0.5)) +  
geom_errorbar(position = position_dodge(width = 0.5)) +  
ylab("Concentration") +  
ggtitle("Variability between wells among sites by contaminant (no calcium)") +  
coord_flip() +  
theme(legend.position="bottom")
```

```
## Warning: Removed 20 rows containing missing values (geom_errorbar).
```

Variability between wells among sites by contaminant (no calcium)

